




# How to improve your fuzzer

- Dictionaries
- Seed corpus
- Custom options
- Fuzzing of non-raw data arguments
- Optimization





# Dictionaries





# Dictionaries

## XML:

```
" encoding=\"1\""  
" a=\"1\""  
" href=\"1\""  
" standalone=\"no\""  
" version=\"1\""  
" xml:base=\"1\""  
" xml:id=\"1\""  
" xml:lang=\"1\""  
" xml:space=\"1\""  
" xmlns=\"1\""  
"&lt;"  
"&#1;"  
"&a;"  
"&#x1;"  
"ANY"  
"[]"  
"CDATA"  
":fallback"  
":include"  
"--"  
"EMPTY"  
"\\\""  
"'''
```

## PNG:

```
header_png="\x89PNG\x0d\x0a\x1a\x0a"  
"  
section_IDAT="IDAT"  
section_IEND="IEND"  
section_IHDR="IHDR"  
section_PLTE="PLTE"  
section_bKGD="bKGD"  
section_cHRM="cHRM"  
section_fRAc="fRAc"  
section_gAMA="gAMA"  
section_gIFg="gIFg"  
section_gIFt="gIFt"  
section_gIFx="gIFx"  
section_hIST="hIST"  
section_iCCP="iCCP"  
section_iTXt="iTXt"  
section_oFFs="oFFs"  
section_pCAL="pCAL"  
section_pHYs="pHYs"  
section_sBIT="sBIT"  
section_sCAL="sCAL"  
section_sPLT="sPLT"  
section_sRGB="sRGB"
```





# Dictionary generation

| Fuzzer                              | Dict. size | Corpus, %     | Coverage, %   | Speed, % |
|-------------------------------------|------------|---------------|---------------|----------|
| expat_xml_parse_fuzzer              | 428        | -3.68         | -2.72         | 1.59     |
| libxml_xml_read_memory_fuzzer       | 780        | 30.81         | 41.14         | 24.13    |
| net_ftp_directory_listing_fuzzer    | 842        | 1.84          | 1.03          | -16.02   |
| net_http_proxy_client_socket_fuzzer | 2,006      | <b>303.90</b> | 26.36         | -0.74    |
| net_http_stream_parser_fuzzer       | 2,006      | <b>259.34</b> | 31.16         | 3.96     |
| net_url_request_fuzzer              | 1,239      | <b>127.44</b> | 6.78          | 4.86     |
| net_websocket_frame_parser_fuzzer   | 925        | 3.85          | 0.35          | 0.77     |
| pdf_css_fuzzer                      | 2,037      | -7.56         | -0.22         | -1.79    |
| pdf_xml_fuzzer                      | 927        | 8.04          | 0.74          | 4.52     |
| sqlite3_prepare_v2_fuzzer           | 657        | <b>370.04</b> | <b>201.68</b> | -11.21   |
| url_parse_fuzzer                    | 650        | <b>473.33</b> | <b>325.07</b> | -11.70   |
| v8_script_parser_fuzzer             | 1,535      | 25.63         | 2.31          | 9.27     |

```
$ chromium/src/testing/libfuzzer/dictionary_generator.py \  
  --fuzzer PATH_TO_FUZZER_BINARY \  
  --spec PATH_TO_FORMAT_SPECIFICATION \  
  --out GENERATED_DICTIONARY.dict
```



# Recommended dictionary

Command: ['<...>/boringssl\_read\_pem\_fuzzer', <...>]

Bot: clusterfuzz-linux-pre-0234

Time ran: 3251.046254

INFO: Seed: 1441206595

#0 READ units: 1166 exec/s: 0

#1167 INITED cov: 204 bits: 366 indir: 16 units: 119 exec/s: 0

#1168 NEW cov: 277 bits: 366 indir: 22 units: 120 exec/s: 0 L: 129

MS: 1 InsertByte-

<...>

#147891067 DONE cov: 278 bits: 370 indir: 22 units: 123 exec/s: 45490

##### Recommended dictionary. #####

"(:\xe7\x9d/!;!;0" # Uses: 8451959

"-----BEGIN " # Uses: 8454876

##### End of recommended dictionary. #####

Done 147891067 runs in 3251 second(s)

stat::number\_of\_executed\_units: 147891067

stat::average\_exec\_per\_sec: 45490

<...>





# Recommended dictionary

Command: ['<...>/net\_url\_request\_fuzzer, <...>']

Bot: clusterfuzz-linux-high-end-pre-0034

Time ran: 3251.478762

Dictionary: 128 entries

INFO: Seed: 964796221

#0 READ units: 10055 exec/s: 0

#128 pulse cov: 5322 bits: 7463 indir: 415 units: 10055 exec/s: 64

<...>

#366810 DONE cov: 12167 bits: 55024 indir: 756 units: 10093 exec/s: 112

##### Recommended dictionary. #####

"blob" # Uses: 3996

"http" # Uses: 2420

"file" # Uses: 1536

"foo" # Uses: 1105

"ntlm" # Uses: 922

##### End of recommended dictionary. #####

<...>





# Recommended dictionary

Command: ['<...>/net\_parse\_cookie\_line\_fuzzer', <...>]

Bot: clusterfuzz-linux-high-end-pre-0080

Time ran: 3251.114852

INFO: Seed: 1831825350

#0 READ units: 1142 exec/s: 0

#1142 INITED cov: 318 bits: 1041 indir: 35 units: 332 exec/s: 0

#1143 NEW cov: 427 bits: 1041 indir: 41 units: 333 exec/s: 0 L: 29

MS: 1 ChangeByte-

<...>

#6808457 NEW cov: 428 bits: 1045 indir: 41 units: 337 exec/s: 11272

L: 698 MS: 5

ChangeByte-AddFromTempAutoDict-ShuffleBytes-AddFromTempAutoDict-CrossOver

- DE: "**httponly**"-"**path**"-

<...>

#36643517 DONE cov: 428 bits: 1045 indir: 41 units: 337 exec/s: 11271

##### Recommended dictionary. #####

**"httponly"** # Uses: 1917525

**"path"** # Uses: 1942187

##### End of recommended dictionary. #####

<...>



# Seed corpus

```
$ ls testing/libfuzzer/fuzzers/woff2_corpus/
```

AhemSpaceLigature.woff2

DejaVuSerif.woff2

MEgalopolisExtra.woff2

Ahem.woff2

EzraSIL.woff2

mplus-1p-regular.woff2

DejaVuSerif-webfont.woff2


LinLibertineO.woff2

OpenSans.woff2

tcu-font.woff2







# Seed corpus

```
$ ls testing/libfuzzer/fuzzers/woff2_corpus/
```

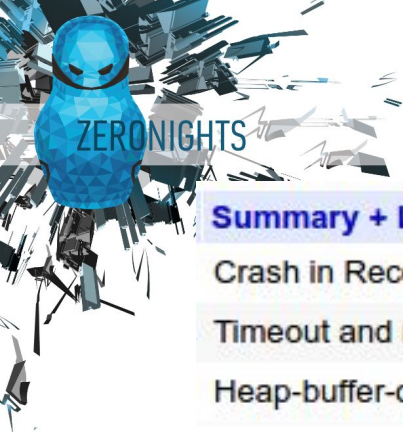
```
AhemSpaceLigature.woff2    DejaVuSerif.woff2    MEgalopolisExtra.woff2  
Ahem.woff2                EzraSIL.woff2        mplus-1p-regular.woff2  
DejaVuSerif-webfont.woff2  LinLibertineO.woff2  OpenSans.woff2  
tcu-font.woff2
```

```
fuzzer_test("convert_woff2ttf_fuzzer") {  
  sources = [  
    "convert_woff2ttf_fuzzer.cc",  
  ]  
  deps = [  
    "//third_party/woff2:woff2_dec",  
  ]  
  seed_corpus = "//testing/libfuzzer/fuzzers/woff2_corpus"  
  libfuzzer_options = [ "max_len=803500" ]  
}
```



# Custom options

```
fuzzer_test("convert_woff2ttf_fuzzer") {  
    sources = [  
        "convert_woff2ttf_fuzzer.cc",  
    ]  
    deps = [  
        "//third_party/woff2:woff2_dec",  
    ]  
    seed_corpus = "//testing/libfuzzer/fuzzers/woff2_corpus"  
    libfuzzer_options = [ "max_len=803500" ]  
}
```



# Custom options

## Summary + Labels ▼

Crash in ReconstructTransformedHmtx Reproducible ClusterFuzz

Timeout and memory leaks in woff2::ConvertWOFF2ToTTF()

Heap-buffer-overflow in Read Reproducible Clusterfuzz

```
fuzzer_test("convert_woff2ttf_fuzzer") {  
  sources = [  
    "convert_woff2ttf_fuzzer.cc",  
  ]  
  deps = [  
    "//third_party/woff2:woff2_dec",  
  ]  
  seed_corpus = "//testing/libfuzzer/fuzzers/woff2_corpus"  
  libfuzzer_options = [ "max_len=803500" ]  
}
```





# Optimization [initial version]

```
#include "third_party/zlib/zlib.h"

// Entry point for LibFuzzer.
extern "C" int LLVMFuzzerTestOneInput(const unsigned char *data, size_t size) {
    uint8_t buffer[1024 * 1024] = { 0 };
    size_t buffer_length = sizeof(buffer);

    if (Z_OK != uncompress(buffer, &buffer_length, data, size)) {
        return 0;
    }

    return 0;
}
```




# Optimization [initial version]

```
#include "third_party/zlib/zlib.h"

// Entry point for LibFuzzer.
extern "C" int LLVMFuzzerTestOneInput(const unsigned char *data, size_t size) {
    uint8_t buffer[1024 * 1024] = { 0 };
    size_t buffer_length = sizeof(buffer);

    if (Z_OK != uncompress(buffer, &buffer_length, data, size)) {
        return 0;
    }

    return 0;
}
```



# Optimization [why you did that?]

```
#include "third_party/zlib/zlib.h"
```

```
// Entry point for LibFuzzer.
```

```
extern "C" int LLVMFuzzerTestOneInput(const unsigned char *data, size_t size) {
```

```
    const int NUM_ITEMS = 1024 * 1024;
```

```
    const int BUF_SIZE = NUM_ITEMS * sizeof(uint8_t);
```

```
    uint8_t *buffer = new uint8_t[NUM_ITEMS];
```

```
    uLongf buffer_length = (uLongf)BUF_SIZE;
```

```
    memset(buffer, 0, BUF_SIZE);
```

```
    if (Z_OK != uncompress(buffer, &buffer_length, data, static_cast<uLong>(size))) {
```

```
        delete[] buffer;
```

```
        return 0;
```

```
    }
```

```
    delete[] buffer;
```

```
    return 0;
```



# Optimization [reanimation]

```
#include "third_party/zlib/zlib.h"
```

```
// Entry point for LibFuzzer.
```

```
extern "C" int LLVMFuzzerTestOneInput(const unsigned char *data, size_t size) {
```

```
    const int NUM_ITEMS = 1024 * 1024;
```

```
    const int BUF_SIZE = NUM_ITEMS * sizeof(uint8_t);
```

```
    uint8_t *buffer = new uint8_t[NUM_ITEMS];
```

```
    uLongf buffer_length = (uLongf)BUF_SIZE;
```

```
    memset(buffer, 0, BUF_SIZE);
```

```
    if (Z_OK != uncompress(buffer, &buffer_length, data, static_cast<uLong>(size))) {
```


```
        delete[] buffer;
```

```
        return 0;
```

```
    }
```

```
    delete[] buffer;
```

```
    return 0;
```




# Optimization [reanimation]

```
#include "third_party/zlib/zlib.h"
static Bytef buffer[256 * 1024] = { 0 };
// Entry point for LibFuzzer.
extern "C" int LLVMFuzzerTestOneInput(const unsigned char *data, size_t size) {
    uLongf buffer_length = static_cast<uLongf>(sizeof(buffer));
    const int BUF_SIZE = NUM_ITEMS * sizeof(uint8_t);
    uint8_t *buffer = new uint8_t[NUM_ITEMS];
    uLongf buffer_length = (uLongf)BUF_SIZE;
    memset(buffer, 0, BUF_SIZE);

    if (Z_OK != uncompress(buffer, &buffer_length, data, static_cast<uLong>(size))) {
        delete[] buffer;
        return 0;
    }

    delete[] buffer;
    return 0;
}
```





# Optimization [final version]

```
#include "third_party/zlib/zlib.h"
```

```
static Bytef buffer[256 * 1024] = { 0 };
```

```
// Entry point for LibFuzzer.
```

```
extern "C" int LLVMFuzzerTestOneInput(const uint8_t* data, size_t size) {
```

```
    uLongf buffer_length = static_cast<uLongf>(sizeof(buffer));
```

```
    if (Z_OK != uncompress(buffer, &buffer_length, data, static_cast<uLong>(size))) {  
        return 0;
```

```
    }
```

```
    return 0;
```

```
}
```



# Guess performance difference?





# Guess performance difference?

- remove memset() for 1MB heap buffer -> ~**5x**





# Guess performance difference?

- remove memset() for 1MB heap buffer -> ~**5x**
- move 1MB heap buffer to the stack -> ~**2x**






# Guess performance difference?

- remove memset() for 1MB heap buffer -> ~**5x**
- move 1MB heap buffer to the stack -> ~**2x**
- use 256KB instead of 1Mb on the stack -> ~**3x**





# Guess performance difference?

- remove memset() for 1MB heap buffer -> ~**5x**
- move 1MB heap buffer to the stack -> ~**2x**
- use 256KB instead of 1Mb on the stack -> ~**3x**
- move 256KB buffer from stack to global -> ~**2x**





# Guess performance difference?

- remove memset() for 1MB heap buffer -> **~5x**
- move 1MB heap buffer to the stack -> **~2x**
- use 256KB instead of 1Mb on the stack -> **~3x**
- move 256KB buffer from stack to global -> **~2x**



**53x**

